

Mudcard

- **How do we actually do the gradient descent function, I know you said not to use the code that you did? Is it a simple sklearn code like it is for linear regression?**
 - Yes, when you call `.fit` of linear or logistic regression, it does gradient descent for you
- **can the learning rate be anything?**
 - If it's too small, it will take a long time to converge.
 - If it's too large, it will introduce numerical instabilities.
 - Learning rates in the range of $1e-4$ and $1e-2$ usually work well.
- **Could you explain the gradient descent a little more?**
 - We need to move on but there are countless online resources to learn about it
- **how are iterations determined while running gradient descent?**
 - do you mean the number of iterations?
 - You can do a couple of things:
 - have a preset number of iterations like in my code
 - have a preset number of operations but stop sooner if a convergence criteria is fulfilled
 - use a while loop and iterate until a convergence criteria is fulfilled (not usually done because the while loop can go forever)

Regularization

By the end of this lecture, you will be able to

- Describe why regularization is important and what are the two types of regularization
- Describe how regularized linear regression works
- Describe how regularized logistic regression works

Regularization

By the end of this lecture, you will be able to

- **Describe why regularization is important and what are the two types of regularization**
- Describe how regularized linear regression works
- Describe how regularized logistic regression works

Polynomial regression


Let's work with a new example dataset

```
In [1]: # load packages
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import PolynomialFeatures
import matplotlib
matplotlib.rcParams.update({'font.size': 11})

df = pd.read_csv('data/regularization_example.csv')
X_ori = df['x0'].values.reshape(-1, 1)
y = df['y'].values
print(np.shape(X_ori))
print(np.shape(y))

# visualize the data
plt.figure(figsize=(5,3))
plt.scatter(X_ori,y)
plt.xlabel('feature')
plt.ylabel('target variable')
plt.show()
```

```
(40, 1)
(40,)
```



```
In [2]: # lets generate more features because a linear model will obviously be insufficient to fit the data well
pf = PolynomialFeatures(degree = 20, include_bias=False)
X = pf.fit_transform(X_ori)
print(np.shape(X))
print(pf.get_feature_names_out())

(40, 20)
['x0' 'x0^2' 'x0^3' 'x0^4' 'x0^5' 'x0^6' 'x0^7' 'x0^8' 'x0^9' 'x0^10'
 'x0^11' 'x0^12' 'x0^13' 'x0^14' 'x0^15' 'x0^16' 'x0^17' 'x0^18' 'x0^19'
 'x0^20']
```

We split data into train and validation!

```
In [3]: from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=1)
print(np.shape(X_train), np.shape(y_train))
print(np.shape(X_val), np.shape(y_val))

(32, 20) (32,)
(8, 20) (8,)
```

Let's train and validate some linear regression models

Use the first feature only

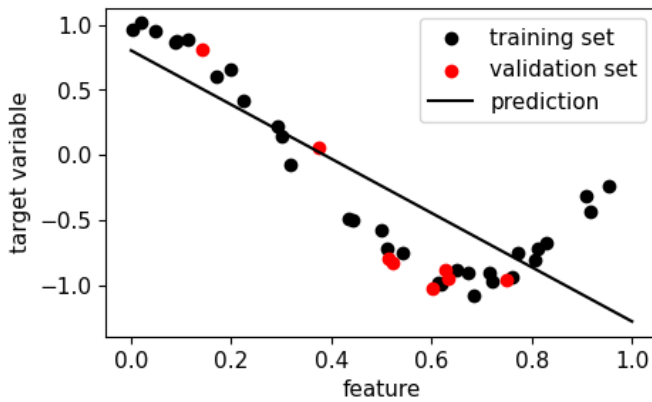
```
In [4]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# let's use only the first feature
linreg = LinearRegression(fit_intercept=True)
linreg.fit(X_train[:, :1], y_train)
print('intercept:', linreg.intercept_)
print('w:', linreg.coef_)

train_MSE = mean_squared_error(y_train, linreg.predict(X_train[:, :1]))
val_MSE = mean_squared_error(y_val, linreg.predict(X_val[:, :1]))
print('train MSE:', train_MSE)
print('val MSE:', val_MSE)

# let's visualize the model
x_model = np.linspace(0, 1, 100)
plt.figure(figsize=(5, 3))
plt.scatter(X_train[:, 0], y_train, color='k', label='training set')
plt.scatter(X_val[:, 0], y_val, color='r', label='validation set')
plt.plot(x_model, linreg.predict(x_model.reshape(-1, 1)), color='k', label='prediction')
plt.xlabel('feature')
plt.ylabel('target variable')
plt.legend()
plt.show()

intercept: 0.8018842867499771
w: [-2.08151827]
train MSE: 0.13964692457239292
val MSE: 0.17142516062337293
```



Use all features


```
In [5]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# use all features
linreg = LinearRegression(fit_intercept=True)
linreg.fit(X_train, y_train)
print('intercept:', linreg.intercept_)
print('ws:', linreg.coef_)

train_MSE = mean_squared_error(y_train, linreg.predict(X_train))
val_MSE = mean_squared_error(y_val, linreg.predict(X_val))
print('train MSE:', train_MSE)
print('val MSE:', val_MSE)

# let's visualize the model
x_model = np.linspace(0,1,1000)
plt.figure(figsize=(5,3))
plt.scatter(X_train[:,0], y_train, color='k', label='training set')
plt.scatter(X_val[:,0], y_val, color='r', label='validation set')
plt.plot(x_model, linreg.predict(pf.transform(x_model.reshape(-1,1))), color='k', label='prediction')
plt.ylim([-1,1])
plt.xlabel('feature')
plt.ylabel('target variable')
plt.legend()
plt.show()
```

```
intercept: 0.9521398936541472
ws: [ 2.94556467e+00  1.78575034e+02 -1.07852992e+04  7.71588159e+04
  3.57083482e+06 -9.66895422e+07  1.20126203e+09 -9.39908662e+09
  5.13303355e+10 -2.05802411e+11  6.23129909e+11 -1.44706005e+12
  2.59415710e+12 -3.58595231e+12  3.78788495e+12 -3.00091345e+12
  1.72536419e+12 -6.79474340e+11  1.63872612e+11 -1.82456642e+10]
train MSE: 0.0022236975806119337
val MSE: 0.032870916690386186
```



What to do?

- the model is visibly performs poorly when only the original feature is used

- the model performs very good on the training set but poorly on the validation set when all features are used
 - the ws are huge!

Regularization solves this problem!

Regularization

By the end of this lecture, you will be able to

- Describe why regularization is important and what are the two types of regularization
- Describe how regularized linear regression works**
- Describe how regularized logistic regression works

Regularization to the rescue!

- let's change the cost function and add a **penalty term** for large ws
- Lasso regression**: regularize using the l1 norm of w:

$$L(w) = \frac{1}{n} \sum_{i=1}^n [(w_0 + \sum_{j=1}^m w_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m |w_j|$$

- Ridge regression**: regularize using the l2 norm of w:

$$L(w) = \frac{1}{n} \sum_{i=1}^n [(w_0 + \sum_{j=1}^m w_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m w_j^2$$

- α is the regularization parameter (positive number), it describes how much we penalize large ws
- With the cost function changed, the derivatives in gradient descent need to be updated too!

Feature selection with Lasso regularization

- Least Absolute Shrinkage and Selection Operator
- cost = MSE + α * l1 norm of w

$$L(w) = \frac{1}{n} \sum_{i=1}^n [(w_0 + \sum_{j=1}^m w_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m |w_j|$$

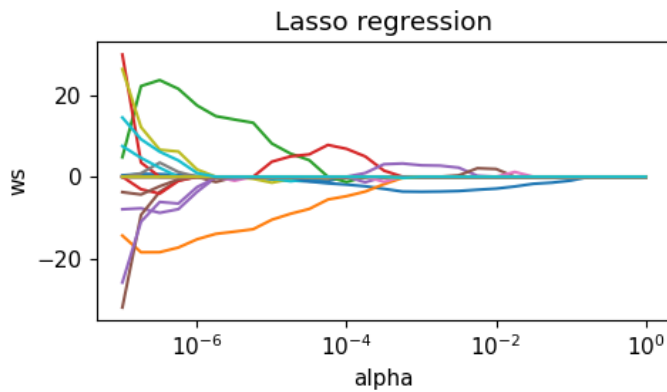
- ideal for feature selection
- as α increases, more and more feature weights are reduced to 0.

```
In [6]: from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error


alpha = np.logspace(-7,0,29)
ws = []
models = []
train_MSE = np.zeros(len(alpha))
val_MSE = np.zeros(len(alpha))

# do the fit
for i in range(len(alpha)):
    # load the linear regression model
    lin_reg = Lasso(alpha=alpha[i],max_iter=100000000)
    lin_reg.fit(X_train, y_train)
    ws.append(lin_reg.coef_)
    models.append(lin_reg)
    train_MSE[i] = mean_squared_error(y_train,lin_reg.predict(X_train))
    val_MSE[i] = mean_squared_error(y_val,lin_reg.predict(X_val))
```

```
In [7]: plt.figure(figsize=(5,3))
plt.plot(alpha, ws)
plt.semilogx()
plt.xlabel('alpha')
plt.ylabel('ws')
plt.title('Lasso regression')
plt.tight_layout()
plt.savefig('figures/lasso_coefs.png',dpi=300)
plt.show()
```



```
In [8]: plt.figure(figsize=(5,3))
plt.plot(alpha,train_MSE,label='train score')
plt.plot(alpha,val_MSE,label='val score')
plt.semilogy()
plt.semilogx()
plt.xlabel('alpha')
plt.ylabel('MSE')
plt.legend()
plt.tight_layout()
plt.savefig('figures/train_val_MSE_lasso.png',dpi=300)
plt.show()
```



Bias vs variance

- Bias: the model performs poorly on both the train and validation sets
 - high alpha in our example
- the model performs very well on the training set but it performs poorly on the validation set
 - low alpha in our example
 - lowering the alpha further would improve the train score but the validation score would increase
 - we don't do it because of convergence issues

The bias-variance trade off

- the curve of the validation score as a function of a hyper-parameter usually has a U shape if evaluation metric needs to be minimized, or an inverted U if the metric needs to be maximized
- choose the hyper-parameter value that gives you the best validation score

Quiz

Which alpha value gives the best validation score? Visualize that model!

In []:

The bias-variance tradeoff with Ridge regularization

- cost = MSE + α * l2 norm of w

$$L(w) = \frac{1}{n} \sum_{i=1}^n [(w_0 + \sum_{j=1}^m w_j x_{ij} - y_i)^2] + \frac{\alpha}{m} \sum_{j=0}^m w_j^2$$

- as α approaches 0, we reproduce the linear regression weights
- small α creates high variance
- large α creates high bias

```
In [9]: from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

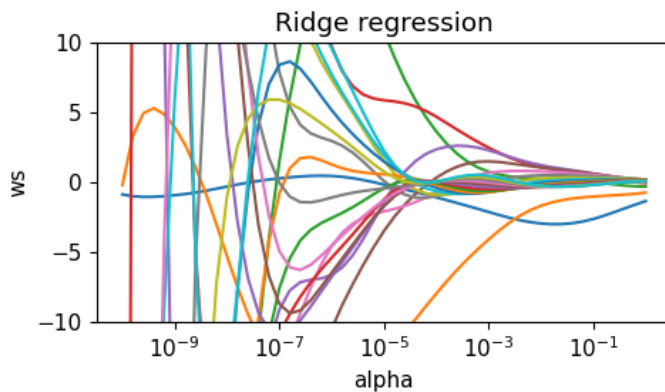
alpha = np.logspace(-10,0,51)

# arrays to save train and test MSE scores
train_MSE = np.zeros(len(alpha))
val_MSE = np.zeros(len(alpha))


ws = []

# do the fit
for i in range(len(alpha)):
    # load the linear regression model
    lin_reg = Ridge(alpha=alpha[i])
    lin_reg.fit(X_train, y_train)
    ws.append(lin_reg.coef_)
    # train and test scores
    train_MSE[i] = mean_squared_error(y_train, lin_reg.predict(X_train))
    val_MSE[i] = mean_squared_error(y_val, lin_reg.predict(X_val))
```

```
In [10]: plt.figure(figsize=(5,3))
plt.plot(alpha, ws)
plt.semilogx()
plt.ylim([-1e1,1e1])
plt.xlabel('alpha')
plt.ylabel('ws')
plt.title('Ridge regression')
plt.tight_layout()
plt.savefig('figures/ridge_coefs.png',dpi=300)
plt.show()
```



```
In [11]: plt.figure(figsize=(5,3))
plt.plot(alpha, train_MSE, label='train score')
plt.plot(alpha, val_MSE, label='val score')
plt.semilogy()
plt.semilogx()
plt.xlabel('alpha')
plt.ylabel('MSE cost function')
plt.legend()
plt.tight_layout()
plt.savefig('figures/train_val_MSE_ridge.png',dpi=300)
plt.show()
```



Quiz

Which α gives us the best tradeoff between bias and variance?

In []:

Regularization

By the end of this lecture, you will be able to

- Describe why regularization is important and what are the two types of regularization
- Describe how regularized linear regression works
- **Describe how regularized logistic regression works**

Logistic regression

- Recap: the logloss metric is the cost function

$$L(w) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln(y'_i) + (1 - y_i) \ln(1 - y'_i)]$$

$$L(w) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln\left(\frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right) + (1 - y_i) \ln\left(1 - \frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right)]$$

- the logloss metric with l1 regularization

$$L(w) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln\left(\frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right) + (1 - y_i) \ln\left(1 - \frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right)] + \frac{\alpha}{m} \sum_{j=0}^m |w_j|$$

- the logloss metric with l2 regularization

$$L(w) = -\frac{1}{N} \sum_{i=1}^n [y_i \ln\left(\frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right) + (1 - y_i) \ln\left(1 - \frac{1}{1 + e^{-w_0 + \sum_{j=1}^m w_j x_{ij}}}\right)] + \frac{\alpha}{m} \sum_{j=0}^m w_j^2$$

Logistic regression in sklearn

```
In [12]: from sklearn.linear_model import LogisticRegression

log_reg_l1 = LogisticRegression(penalty='l1', C = 1/alpha) # C is the inverse of the regularization strength
log_reg_l2 = LogisticRegression(penalty='l2', C = 1/alpha)
# fit, predict, predict_proba are available
# log_reg.coef_ returns the w values
```

```
In [13]: help(LogisticRegression)
```

Help on class LogisticRegression in module sklearn.linear_model._logistic:

```
class LogisticRegression(sklearn.linear_model._base.LinearClassifierMixin, sklearn.linear_model._base.SparseCoefMixin, sklearn.base.BaseEstimator)
| LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the 'multi_class' option is set to 'ovr', and uses the cross-entropy loss if the 'multi_class' option is set to 'multinomial'. (Currently the 'multinomial' option is supported only by the 'lbfgs', 'sag', 'saga' and 'newton-cg' solvers.)

This class implements regularized logistic regression using the 'liblinear' library, 'newton-cg', 'sag', 'saga' and 'lbfgs' solvers. **Note** that regularization is applied by default. It can handle both dense and sparse input. Use C-ordered arrays or CSR matrices containing 64-bit floats for optimal performance; any other input format will be converted (and copied).

The 'newton-cg', 'sag', and 'lbfgs' solvers support only L2 regularization with primal formulation, or no regularization. The 'liblinear' solver supports both L1 and L2 regularization, with a dual formulation only for the L2 penalty. The Elastic-Net regularization is only supported by the 'saga' solver.

Read more in the :ref:`User Guide <logistic_regression>`.

Parameters

penalty : {'l1', 'l2', 'elasticnet', None}, default='l2'

Specify the norm of the penalty:

- 'None': no penalty is added;
- 'l2': add a L2 penalty term and it is the default choice;
- 'l1': add a L1 penalty term;
- 'elasticnet': both L1 and L2 penalty terms are added.

.. warning::

Some penalties may not work with some solvers. See the parameter 'solver' below, to know the compatibility between the penalty and solver.

.. versionadded:: 0.19

l1 penalty with SAGA solver (allowing 'multinomial' + L1)

.. deprecated:: 1.2

The 'none' option was deprecated in version 1.2, and will be removed in 1.4. Use 'None' instead.

dual : bool, default=False

Dual or primal formulation. Dual formulation is only implemented for l2 penalty with liblinear solver. Prefer dual=False when n_samples > n_features.

tol : float, default=1e-4

Tolerance for stopping criteria.

C : float, default=1.0

Inverse of regularization strength; must be a positive float. Like in support vector machines, smaller values specify stronger regularization.

fit_intercept : bool, default=True

Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.

intercept_scaling : float, default=1

Useful only when the solver 'liblinear' is used and self.fit_intercept is set to True. In this case, x becomes [x, self.intercept_scaling], i.e. a "synthetic" feature with constant value equal to intercept_scaling is appended to the instance vector. The intercept becomes 'intercept_scaling * synthetic_feature_weight'.

Note! the synthetic feature weight is subject to l1/l2 regularization


```

as all other features.
To lessen the effect of regularization on synthetic feature weight
(and therefore on the intercept) intercept_scaling has to be increased.

class_weight : dict or 'balanced', default=None
Weights associated with classes in the form ``{class_label: weight}``.
If not given, all classes are supposed to have weight one.

The "balanced" mode uses the values of y to automatically adjust
weights inversely proportional to class frequencies in the input data
as ``n_samples / (n_classes * np.bincount(y))``.

Note that these weights will be multiplied with sample_weight (passed
through the fit method) if sample_weight is specified.

.. versionadded:: 0.17
   *class_weight='balanced'*

random_state : int, RandomState instance, default=None
Used when ``solver`` == 'sag', 'saga' or 'liblinear' to shuffle the
data. See :term:`Glossary <random_state>` for details.

solver : {'lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'}, default='lbfgs'

Algorithm to use in the optimization problem. Default is 'lbfgs'.
To choose a solver, you might want to consider the following aspects:

- For small datasets, 'liblinear' is a good choice, whereas 'sag'
  and 'saga' are faster for large ones;
- For multiclass problems, only 'newton-cg', 'sag', 'saga' and
  'lbfgs' handle multinomial loss;
- 'liblinear' is limited to one-versus-rest schemes.
- 'newton-cholesky' is a good choice for ``n_samples`` >> ``n_features``,
  especially with one-hot encoded categorical features with rare
  categories. Note that it is limited to binary classification and the
  one-versus-rest reduction for multiclass classification. Be aware that
  the memory usage of this solver has a quadratic dependency on
  ``n_features`` because it explicitly computes the Hessian matrix.

.. warning::
The choice of the algorithm depends on the penalty chosen.
Supported penalties by solver:

- 'lbfgs'           - ['l2', None]
- 'liblinear'       - ['l1', 'l2']
- 'newton-cg'        - ['l2', None]
- 'newton-cholesky' - ['l2', None]
- 'sag'              - ['l2', None]
- 'saga'             - ['elasticnet', 'l1', 'l2', None]

.. note::
'sag' and 'saga' fast convergence is only guaranteed on features
with approximately the same scale. You can preprocess the data with
a scaler from :mod:`sklearn.preprocessing`.

.. seealso::
Refer to the User Guide for more information regarding
:class:`LogisticRegression` and more specifically the
:ref:`Table <Logistic_regression>`
summarizing solver/penalty supports.

.. versionadded:: 0.17
Stochastic Average Gradient descent solver.
.. versionadded:: 0.19
SAGA solver.
.. versionchanged:: 0.22
The default solver changed from 'liblinear' to 'lbfgs' in 0.22.
.. versionadded:: 1.2
newton-cholesky solver.

max_iter : int, default=100
Maximum number of iterations taken for the solvers to converge.

multi_class : {'auto', 'ovr', 'multinomial'}, default='auto'
If the option chosen is 'ovr', then a binary problem is fit for each
label. For 'multinomial' the loss minimised is the multinomial loss fit
across the entire probability distribution, *even* when the data is
binary*. 'multinomial' is unavailable when solver='liblinear'.
'auto' selects 'ovr' if the data is binary, or if solver='liblinear',

```

```

and otherwise selects 'multinomial'.

.. versionadded:: 0.18
    Stochastic Average Gradient descent solver for 'multinomial' case.
.. versionchanged:: 0.22
    Default changed from 'ovr' to 'auto' in 0.22.

verbose : int, default=0
    For the liblinear and lbfgs solvers set verbose to any positive
    number for verbosity.

warm_start : bool, default=False
    When set to True, reuse the solution of the previous call to fit as
    initialization, otherwise, just erase the previous solution.
    Useless for liblinear solver. See :term:`the Glossary <warm_start>`.

.. versionadded:: 0.17
    *warm_start* to support *lbfgs*, *newton-cg*, *sag*, *saga* solvers.

n_jobs : int, default=None
    Number of CPU cores used when parallelizing over classes if
    multi_class='ovr'. This parameter is ignored when the ``solver`` is
    set to 'liblinear' regardless of whether 'multi_class' is specified or
    not. ``None`` means 1 unless in a :obj:`joblib.parallel_backend`
    context. ``-1`` means using all processors.
    See :term:`Glossary <n_jobs>` for more details.

l1_ratio : float, default=None
    The Elastic-Net mixing parameter, with ``0 <= l1_ratio <= 1``. Only
    used if ``penalty='elasticnet'``. Setting ``l1_ratio=0`` is equivalent
    to using ``penalty='l2'``, while setting ``l1_ratio=1`` is equivalent
    to using ``penalty='l1'``. For ``0 < l1_ratio < 1``, the penalty is a
    combination of L1 and L2.

```

Attributes

```

classes_ : ndarray of shape (n_classes, )
    A list of class labels known to the classifier.

coef_ : ndarray of shape (1, n_features) or (n_classes, n_features)
    Coefficient of the features in the decision function.

    `coef_` is of shape (1, n_features) when the given problem is binary.
    In particular, when `multi_class='multinomial'`, `coef_` corresponds
    to outcome 1 (True) and `-coef_` corresponds to outcome 0 (False).

intercept_ : ndarray of shape (1,) or (n_classes,)
    Intercept (a.k.a. bias) added to the decision function.

    If `fit_intercept` is set to False, the intercept is set to zero.
    `intercept_` is of shape (1,) when the given problem is binary.
    In particular, when `multi_class='multinomial'`, `intercept_`
    corresponds to outcome 1 (True) and `-intercept_` corresponds to
    outcome 0 (False).

n_features_in_ : int
    Number of features seen during :term:`fit`.

.. versionadded:: 0.24

feature_names_in_ : ndarray of shape (`n_features_in_`,)
    Names of features seen during :term:`fit`. Defined only when `X`
    has feature names that are all strings.

.. versionadded:: 1.0

n_iter_ : ndarray of shape (n_classes,) or (1, )
    Actual number of iterations for all classes. If binary or multinomial,
    it returns only 1 element. For liblinear solver, only the maximum
    number of iteration across all classes is given.

.. versionchanged:: 0.20

    In SciPy <= 1.0.0 the number of lbfgs iterations may exceed
    ``max_iter``. ``n_iter_`` will now report at most ``max_iter``.

```

See Also

```

SGDClassifier : Incrementally trained logistic regression (when given
the parameter ``loss="log_loss"``).
LogisticRegressionCV : Logistic regression with built-in cross validation.

Notes
-----
The underlying C implementation uses a random number generator to
select features when fitting the model. It is thus not uncommon,
to have slightly different results for the same input data. If
that happens, try with a smaller tol parameter.

Predict output may not match that of standalone liblinear in certain
cases. See :ref:`differences from liblinear <liblinear_differences>`
in the narrative documentation.

References
-----
L-BFGS-B -- Software for Large-scale Bound-constrained Optimization
Ciyou Zhu, Richard Byrd, Jorge Nocedal and Jose Luis Morales.
http://users.iems.northwestern.edu/~nocedal/lbfgsb.html

LIBLINEAR -- A Library for Large Linear Classification
https://www.csie.ntu.edu.tw/~cjlin/liblinear/

SAG -- Mark Schmidt, Nicolas Le Roux, and Francis Bach
Minimizing Finite Sums with the Stochastic Average Gradient
https://hal.inria.fr/hal-00860051/document

SAGA -- Defazio, A., Bach F. & Lacoste-Julien S. (2014).
:arxiv:"SAGA: A Fast Incremental Gradient Method With Support
for Non-Strongly Convex Composite Objectives" <1407.0202>`

Hsiang-Fu Yu, Fang-Lan Huang, Chih-Jen Lin (2011). Dual coordinate descent
methods for logistic regression and maximum entropy models.
Machine Learning 85(1-2):41-75.
https://www.csie.ntu.edu.tw/~cjlin/papers/maxent\_dual.pdf

Examples
-----
>>> from sklearn.datasets import load_iris
>>> from sklearn.linear_model import LogisticRegression
>>> X, y = load_iris(return_X_y=True)
>>> clf = LogisticRegression(random_state=0).fit(X, y)
>>> clf.predict(X[:2, :])
array([0, 0])
>>> clf.predict_proba(X[:2, :])
array([[9.8...e-01, 1.8...e-02, 1.4...e-08],
       [9.7...e-01, 2.8...e-02, ...e-08]])
>>> clf.score(X, y)
0.97...

Method resolution order:
LogisticRegression
sklearn.linear_model._base.LinearClassifierMixin
sklearn.base.ClassifierMixin
sklearn.linear_model._base.SparseCoefMixin
sklearn.base.BaseEstimator
sklearn.utils.metadata_requests._MetadataRequester
builtins.object

Methods defined here:

__init__(self, penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
    Initialize self. See help(type(self)) for accurate signature.

fit(self, X, y, sample_weight=None)
    Fit the model according to the given training data.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    Training vector, where `n_samples` is the number of samples and
    `n_features` is the number of features.

y : array-like of shape (n_samples,)
    Target vector relative to X.

```

```

sample_weight : array-like of shape (n_samples,) default=None
    Array of weights that are assigned to individual samples.
    If not provided, then each sample is given unit weight.

    .. versionadded:: 0.17
       *sample_weight* support to LogisticRegression.

Returns
-----
self
    Fitted estimator.

Notes
-----
The SAGA solver supports both float64 and float32 bit arrays.

predict_log_proba(self, X)
    Predict logarithm of probability estimates.

    The returned estimates for all classes are ordered by the
    label of classes.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Vector to be scored, where `n_samples` is the number of samples and
    `n_features` is the number of features.

Returns
-----
T : array-like of shape (n_samples, n_classes)
    Returns the log-probability of the sample for each class in the
    model, where classes are ordered as they are in ``self.classes_``.

predict_proba(self, X)
    Probability estimates.

    The returned estimates for all classes are ordered by the
    label of classes.

    For a multi-class problem, if multi_class is set to be "multinomial"
    the softmax function is used to find the predicted probability of
    each class.
    Else use a one-vs-rest approach, i.e calculate the probability
    of each class assuming it to be positive using the logistic function.
    and normalize these values across all the classes.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Vector to be scored, where `n_samples` is the number of samples and
    `n_features` is the number of features.

Returns
-----
T : array-like of shape (n_samples, n_classes)
    Returns the probability of the sample for each class in the model,
    where classes are ordered as they are in ``self.classes_``.

set_fit_request(self: sklearn.linear_model._logistic.LogisticRegression, *, sample_weight: Union[bool, NoneType]
e, str) = '$UNCHANGED$') -> sklearn.linear_model._logistic.LogisticRegression
    Request metadata passed to the ``fit`` method.

    Note that this method is only relevant if
    ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
    Please see :ref:`User Guide <metadata_routing>` on how the routing
    mechanism works.

    The options for each parameter are:

    - ``True``: metadata is requested, and passed to ``fit`` if provided. The request is ignored if metadata i
s not provided.

    - ``False``: metadata is not requested and the meta-estimator will not pass it to ``fit``.

    - ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.

    - ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original n

```

```

ame.
|
| The default ('`sklearn.utils.metadata_routing.UNCHANGED`') retains the
| existing request. This allows you to change the request for some
| parameters and not others.
|
| .. versionadded:: 1.3
|
| .. note::
|     This method is only relevant if this estimator is used as a
|     sub-estimator of a meta-estimator, e.g. used inside a
|     :class:`pipeline.Pipeline`. Otherwise it has no effect.
|
| Parameters
| -----
| sample_weight : str, True, False, or None,                default=sklearn.utils.metadata_routing.UNCH
ANGED
|     Metadata routing for ``sample_weight`` parameter in ``fit``.
|
| Returns
| -----
| self : object
|     The updated object.
|
| set_score_request(self: sklearn.linear_model._logistic.LogisticRegression, *, sample_weight: Union[bool, NoneT
ype, str] = '$UNCHANGED$') -> sklearn.linear_model._logistic.LogisticRegression
|     Request metadata passed to the ``score`` method.
|
| Note that this method is only relevant if
| ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
| Please see :ref:`User Guide <metadata_routing>` on how the routing
| mechanism works.
|
| The options for each parameter are:
|
| - ``True``: metadata is requested, and passed to ``score`` if provided. The request is ignored if metadata
is not provided.
|
| - ``False``: metadata is not requested and the meta-estimator will not pass it to ``score``.
|
| - ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
|
| - ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original n
ame.
|
| The default ('`sklearn.utils.metadata_routing.UNCHANGED`') retains the
| existing request. This allows you to change the request for some
| parameters and not others.
|
| .. versionadded:: 1.3
|
| .. note::
|     This method is only relevant if this estimator is used as a
|     sub-estimator of a meta-estimator, e.g. used inside a
|     :class:`pipeline.Pipeline`. Otherwise it has no effect.
|
| Parameters
| -----
| sample_weight : str, True, False, or None,                default=sklearn.utils.metadata_routing.UNCH
ANGED
|     Metadata routing for ``sample_weight`` parameter in ``score``.
|
| Returns
| -----
| self : object
|     The updated object.
|
| -----
| Data and other attributes defined here:
|
| __annotations__ = {'_parameter_constraints': <class 'dict'>}
|
| -----
| Methods inherited from sklearn.linear_model._base.LinearClassifierMixin:
|
| decision_function(self, X)
|     Predict confidence scores for samples.
|
|     The confidence score for a sample is proportional to the signed

```

distance of that sample to the hyperplane.

Parameters

X : {array-like, sparse matrix} of shape (n_samples, n_features)
The data matrix for which we want to get the confidence scores.

Returns

scores : ndarray of shape (n_samples,) or (n_samples, n_classes)
Confidence scores per `(n_samples, n_classes)` combination. In the binary case, confidence score for `self.classes_[1]` where >0 means this class would be predicted.

predict(self, X)
Predict class labels for samples in X.

Parameters

X : {array-like, sparse matrix} of shape (n_samples, n_features)
The data matrix for which we want to get the predictions.

Returns

y_pred : ndarray of shape (n_samples,)
Vector containing the class labels for each sample.

Methods inherited from sklearn.base.ClassifierMixin:

score(self, X, y, sample_weight=None)
Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters

X : array-like of shape (n_samples, n_features)
Test samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs)
True labels for `X`.

sample_weight : array-like of shape (n_samples,), default=None
Sample weights.

Returns

score : float
Mean accuracy of ``self.predict(X)`` w.r.t. `y`.

Data descriptors inherited from sklearn.base.ClassifierMixin:

__dict__
dictionary for instance variables (if defined)

__weakref__
list of weak references to the object (if defined)

Methods inherited from sklearn.linear_model._base.SparseCoefMixin:

densify(self)
Convert coefficient matrix to dense array format.

Converts the ``coef_`` member (back) to a numpy.ndarray. This is the default format of ``coef_`` and is required for fitting, so calling this method is only required on models that have previously been sparsified; otherwise, it is a no-op.

Returns

self
Fitted estimator.

sparsify(self)

Convert coefficient matrix to sparse format.

Converts the ``coef_`` member to a `scipy.sparse` matrix, which for L1-regularized models can be much more memory- and storage-efficient than the usual `numpy.ndarray` representation.

The ``intercept_`` member is not converted.

Returns

self
 Fitted estimator.

Notes

For non-sparse models, i.e. when there are not many zeros in ``coef_``, this may actually *increase* memory usage, so use this method with care. A rule of thumb is that the number of zero elements, which can be computed with ``(coef_ == 0).sum()``, must be more than 50% for this to provide significant benefits.

After calling this method, further fitting with the `partial_fit` method (if any) will not work until you call `densify`.

Methods inherited from `sklearn.base.BaseEstimator`:

`__getstate__(self)`
 Helper for pickle.

`__repr__(self, N_CHAR_MAX=700)`
 Return `repr(self)`.

`__setstate__(self, state)`

`__sklearn_clone__(self)`

`get_params(self, deep=True)`
 Get parameters for this estimator.

Parameters

deep : bool, default=True
 If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns

params : dict
 Parameter names mapped to their values.

`set_params(self, **params)`
 Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as `:class:`~sklearn.pipeline.Pipeline``). The latter have parameters of the form ``<component>__<parameter>`` so that it's possible to update each component of a nested object.

Parameters

**params : dict
 Estimator parameters.

Returns

self : estimator instance
 Estimator instance.

Methods inherited from `sklearn.utils._metadata_requests._MetadataRequester`:

`get_metadata_routing(self)`
 Get metadata routing of this object.

Please check `:ref:`User Guide <metadata_routing>`` on how the routing mechanism works.

Returns

```
-----
routing : MetadataRequest
    A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
        routing information.
```

```
-----
Class methods inherited from sklearn.utils._metadata_requests._MetadataRequester:
```

```
__init_subclass__(**kwargs) from builtins.type
    Set the ``set_{method}_request`` methods.
```

This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It looks for the information available in the set default values which are set using ``__metadata_request__*`` class attributes, or inferred from method signatures.

The ``__metadata_request__*`` class attributes are used when a method does not explicitly accept a metadata through its arguments or if the developer would like to specify a request value for those metadata which are different from the default ``None``.

References

```
-----
.. [1] https://www.python.org/dev/peps/pep-0487
```

Mudcard

In []: